

AN INDEPENDENT TUTORIAL · LLM APPLICATION EVALUATION

Evaluating RAG, Agents and Tools

How to evaluate retrieval-augmented generation, tool-using agents, and conversational systems — end to end: what to measure, how to measure it, how to keep the measurement honest, and how to run it in production. Every concept is anchored in one real, runnable example: a customer-support agent evaluated with Databricks' MLflow 3 GenAI stack.

-
- The vocabulary of evaluation — metrics, scorers, judges, traces, gates
 - Retrieval (relevance) vs. generation (correctness, groundedness) metrics
 - Tool behaviour: call correctness, selection, escalation, consent, leakage
 - Custom scorers and LLM judges — contracts, mechanics, and traps
 - Datasets, regression testing, prompt versioning, and promotion gates
 - Online monitoring, judge-human alignment, and automated prompt optimisation

September 2026 · 16 chapters · grounded in the TelcoAssist evaluation workspace

Contents

Sixteen chapters, one thread: decide what “good” means, make it measurable, and gate every change on it — whether a human or a machine wrote the prompt.

01	Why evaluate agents at all	2
02	The system under test	4
03	The vocabulary of evaluation	8
04	Decide what “good” means before you measure	10
05	Traces — the raw material you actually measure	12
06	Retrieval & generation metrics	13
07	Tool behaviour & safety metrics	15
08	Custom scorers & LLM judges	16
09	Building evaluation datasets	18
10	Regression testing & prompt versioning	19
11	Online evaluation & monitoring	21
12	Aligning LLM judges with humans	22
13	Automated prompt optimisation	23
14	Multi-turn conversation evaluation	25
15	The traps that bite everyone	26
16	Reference — the numbers of the track	28

01 Why evaluate agents at all

AGENT BRIEF — HIGH LEVEL

TelcoAssist answers customer questions from an internal knowledge base of ten articles, optionally calls one of three tools (`lookup_account` , `check_network_status` , `open_ticket`) and replies in plain text — single-turn via `answer` , multi-turn via `converse` . That is the whole system under test in this tutorial; everything you learn here is demonstrated against it.

Here is the uncomfortable truth that starts this whole discipline: **evaluating an LLM is not the same as evaluating an LLM application**. A raw model gets judged on the fluency or factual accuracy of its output. Your agent is a *system* — retrieval, a model, tools, memory, prompts, and judgment logic composed together. Each component can be individually fine while the system fails, and the failures that matter (acts without consent, leaks another customer's data, invents a price) are *system* failures, not model failures.

Failures compound — the arithmetic of a pipeline

Assume every step of a request is independently correct 95% of the time. For a single-turn RAG answer there may be two steps (retrieve, then answer) — 90% end-to-end. Give the agent tools and multi-turn memory and you have a chain of model decisions per conversation: retrieving, calling a tool or not, choosing the right tool, refusing an unsafe request, answering faithfully. If each decision is 95% reliable, a five-turn conversation succeeds about **77%** of the time (0.95^5), and by turn fourteen it is below **50%** ($0.95^{14} \approx 49\%$). The agent degrades invisibly, turn by turn, while every per-call metric looks healthy.

Evaluation is not a QA step at the end. It is the only way to know what the composed system actually does — wrong tool calls, invented figures, unconsented writes — before customers find out for you.

Online discovers; offline prevents recurrence

The discipline has two halves that answer two different questions, and a healthy setup runs both:

- **Online evaluation / monitoring** watches production traffic and answers “*what is happening in the wild?*” — novel questions, drift, edge cases you never designed for. Its virtue: it discovers things you did not anticipate. Its weakness: it acts once, in the moment.
- **Offline evaluation** runs a fixed test suite against a fixed dataset with fixed thresholds, before anything ships. Its virtue: it repeats — the same test every release. Its weakness: it only knows what you thought to put in the dataset.

The loop that makes evaluation strong is moving findings from one half to the other: *an interesting trace spotted in production gets mined into the offline dataset, so the next release is protected against that exact failure*. Discover once online, prevent forever offline.

The gate is the only thing that can deploy

Every candidate — a hand-edited prompt, a new model, a prompt an optimiser generated — goes through the *same* quality gate before it reaches production. That single rule removes the politics and the guesswork from releases: the metric says yes or no, the reviewer reads the evidence, and the gate decides. You will see later that even automated prompt optimisation (chapter 13) does not bypass this; it feeds the same gate.

What evaluation actually buys you

- **Decisions you can defend:** “relevance dropped 0.06, gate says no” beats “feels a bit worse”.
- **Regression protection:** a prompt edit that improves one metric often breaks another — the gate catches both.
- **A signal for improvement:** you cannot optimise what you cannot measure. Prompt optimisation in chapter 13 takes the same metrics as its objective.
- **Evidence for stakeholders:** safety and consent gates at 1.00 are the difference between “we tested for it” and “we believe in it”.

THE PRINCIPLE BEHIND EVERYTHING THAT FOLLOWS

Design the system to be measured from the start. The TelcoAssist agent was deliberately built so that retrieval always runs (one scorer per request works), the write tool is ungated at the API level (obedience has to be *verified*, not assumed), and the prompt is a parameter (versioning and optimisation become possible). Every later chapter relies on those three choices.

[Next: the system under test →](#) **Section 01 of 16**

02 The system under test


AGENT BRIEF — HIGH LEVEL

TelcoAssist is a customer-support agent for a fictional telecom company. A customer asks a question; the agent answers from an internal knowledge base of ten support articles, optionally calling one of three tools — `lookup_account` (account data), `check_network_status` (outages), and `open_ticket` (a write action that requires explicit customer consent) — and replies in plain text. It works both single-turn (`answer`) and multi-turn (`converse`). It is deliberately the least interesting part of this tutorial: it exists to be *measurable*, not impressive.

Before any evaluation makes sense, you have to know what an agent is *for* and what it can touch. Everything below is the working definition the whole tutorial evaluates against — functionality, tools, data, models, and the trace it produces.

What it does, end to end

A request flows through a fixed pipeline with one model-driven loop. Retrieval always runs first; whether a tool gets called is the model's decision:



The request flow: retrieve → think → maybe call a tool → answer.

Three design points exist purely so evaluation is possible:

- **Retrieval is a fixed pipeline step, not a tool the model may skip.** That guarantees every trace carries a `RETRIEVER` span, so groundedness is always scoreable. Account lookup *stays* an LLM-chosen tool precisely because “should it have called that?” is itself something later chapters score.
- **check_network_status is functionally trivial and strategically essential.** With one tool, “did it call a tool” and “did it call the *right* tool” are the same question. A second read-only tool separates them — and wrong-tool turns out to be a nastier failure than no-tool, because it returns data that is real, irrelevant, and confidently presented.
- **open_ticket is completely ungated at the tool layer.** Call it and it writes, every time. That is realistic — most agent write-actions are gated by *instructions*, not by the API — and it makes obedience something evaluation has to verify rather than something the code guarantees.

Two entry points

Function	Signature	Returns	Used by
<code>answer</code>	<code>(query, customer_id, system_prompt, history)</code>	<code>str</code>	Single-turn evaluation (chapters 5–11)
<code>converse</code>	<code>(turns, customer_id, system_prompt)</code>	<code>{"response", "turns"}</code>	Multi-turn evaluation (chapter 14)

`converse` threads history between turns and applies `customer_id` on *every* turn — identity is session-scoped. The `response` key exists so every single-turn scorer keeps working unchanged on conversation runs.

Tools

Name	Span type	Read/write	When it should fire
<code>retrieve_kb</code>	<code>RETRIEVER</code>	read	Always — fixed first step, not model-chosen
<code>lookup_account</code>	<code>TOOL</code>	read	Only when the answer needs <i>this</i> customer's own account data and a customer ID is present
<code>check_network_status</code>	<code>TOOL</code>	read	Only for signal/outage questions and an area code is present
<code>open_ticket</code>	<code>TOOL</code>	WRITE	Only after the customer explicitly agreed, in their own words

Each tool is a thin `@tool` wrapper over a separately traced `_impl`, so the TOOL span comes from our own instrumentation rather than from whatever the framework's autologging emits.

Data the agent holds

All in-memory Python literals — no database, no external service. Three read paths and exactly one write path (tickets), and that asymmetry is the whole reason the consent/approval work exists in later chapters.

Store	Contents	Notes
<code>KNOWLEDGE_BASE</code>	10 articles — plans, billing, throttling, late fees, roaming, upgrades, troubleshooting, refunds, account changes, AutoPay	Deliberately fact-dense : concrete figures (\$45/25GB, \$65/75GB, \$85 uncapped; \$15 late fee; \$12/day roaming...) give scorers something specific to assert and make hallucinations obvious. Embedded once, L2-normalised; retrieval is a NumPy dot product over top-3.
<code>ACCOUNTS</code>	3 records: CUST-1001 (Plus, \$78.40, active, AutoPay), CUST-1002 (Essential, \$0.00, active), CUST-1003 (Unlimited, \$152.90, suspended)	Each carries a different combination so tests can distinguish them. <code>CUST-9999</code> is used in the edge cases precisely because it is <i>absent</i> .
<code>NETWORK_STATUS</code>	3 area codes: 415 degraded (6h), 212 operational, 512 outage (18h)	One of each state. Area code <code>999</code> deliberately absent, same reason as <code>CUST-9999</code> .
<code>TICKETS</code>	starts empty; written by <code>open_ticket</code>	The agent's sole side effect and the reason the approval gate exists.

Models

Selected by the `TELCOASSIST_PROVIDER` environment variable; both are created *lazily*, so `import agent` requires no credentials. The model under test is pinned and visible — a silently platform-dependent model would make before/after comparison meaningless.


Provider	Chat model	Embedding model
<code>databricks</code> (default)	<code>databricks-gpt-oss-120b</code>	<code>databricks-gte-large-en</code>
<code>openai</code>	<code>gpt-4o-mini</code>	<code>text-embedding-3-small</code>

The system prompt is a parameter

The prompt is carried in graph state (`answer(..., system_prompt=...)`). That single design choice is what lets regression testing evaluate a version pulled from the MLflow Prompt Registry, and lets an optimiser swap candidates, without touching the agent code. Its clauses map one-to-one onto scorers — the prompt and the eval criteria are written as a matched pair:

Prompt clause	Scored by
Answer only from the support articles	groundedness, abstains-when-unsupported
Call each tool only in its stated situation	tool call correctness, tool selection
Never disclose another customer's data	no-account-leakage, protects-other-accounts
Refunds/changes/cancellations go to a human	escalation
Only open a ticket after explicit agreement	approval before write
Keep answers concise	conciseness, response word count

The trace it produces — evaluation reads this, not just the string



Trace shape: conversation → turn → (retrieval, model, tools). The nesting is load-bearing.

The nesting is what lets the approval scorer work out *which turn* a write landed on and compare it against which turn the customer consented on. Flatten the trace and that scorer becomes impossible to write.

What it deliberately cannot do

Non-goal	Status
Authentication / identity verification	Still excluded — a <code>customer_id</code> in the request is trusted; that itself is what the adversarial suite probes.
Autonomous escalation routing	Still excluded — it says “a human will handle this”; it doesn't route.
Write actions	One , consent-gated: <code>open_ticket</code> .

Multi-turn memory

`converse` threads history — that is the multi-turn chapter's subject.

Those boundaries exist because they constrain the *evaluation*, not just the agent: with one tool, tool-selection accuracy was unmeasurable, and without conversations, “acted without asking” could not even be expressed.

THE PRINCIPLE BEHIND THE WHOLE DESIGN

*“A non-goal that limits the agent is discipline. A non-goal that silently limits what you can **measure** is a blind spot.”* It took the evaluation machinery of the later chapters to make that visible. Keep it in mind whenever you scope an agent you plan to evaluate.

Next: the vocabulary of evaluation → Section 02 of 16

03 The vocabulary of evaluation

AGENT BRIEF — HIGH LEVEL

TelcoAssist answers from the knowledge base, chooses whether to call a tool, and replies in text. Its output is not a single string: the evaluation reads *the answer*, *the tools that ran*, *the retrieved articles*, and *the full trace* — and judges each aspect separately. That is why the vocabulary below has a word for each layer of what it produces.

Every evaluation you will read about is built from a small fixed set of pieces. If you master these terms, the rest of the tutorial is just putting them together.

Term	What it is	Example from this tutorial
Metric	A named, typed measurement taken per evaluation example	relevance, groundedness, tool call correctness
Scorer	The function that computes one metric for one example — turns inputs into a score and (usually) free-text feedback	<code>RetrievalGroundedness</code> , <code>approval_before_write</code>
Judge	An LLM used as a scorer: it reads candidate output(s), applies a rubric, and emits a verdict	the correctness / escalation / tool-selection judges
Deterministic scorer	Runs in Python — same input, same output, zero variance	word-count, response length, exact-match
LLM-judged scorer	A model call — inherently noisy between runs	all semantic metrics here
Dataset	The fixed set of evaluation examples (inputs, expected outputs, ground truth)	the curated TelcoAssist eval dataset

Trace / span	The recorded execution of one agent run; a span is one step (retrieval, model call, tool call)	RETRIEVER , CHAT_MODEL , TOOL , AGENT spans
Aggregation	How per-example scores are summarised into one number for a run	mean, median, p90 — the full six are in the reference
Threshold	The pass/fail level a metric must meet	groundedness ≥ 0.90
Tolerance	How much worse a candidate may be vs. baseline before a regression counts	0.00 deterministic / 0.02 judged
Quality gate	A named, versioned, blocking rule built from metric + threshold	the eleven gates of this track (see chapter 4)

Deterministic vs. judged — the single most consequential split

Everything downstream inherits one property: **deterministic scorers are stable, judged scorers are noisy**. The consequences:

- **Regression tolerance.** If a metric wobbles by ± 0.02 between identical runs purely because a judge is a model, you cannot treat a 0.01 drop as a real regression. Hence the convention: deterministic metrics get a **TOLERANCE 0.00** , judged metrics get **TOLERANCE 0.02** .
- **Comparison design.** To tell two prompts apart you need the judge's noise to be smaller than the effect you are looking for — which is where datasets and sample sizes come in (chapters 9 and 11).
- **Judge quality becomes its own metric.** Because a judge's verdict is load-bearing, you need a process to check the judge itself (chapter 12).

Inputs and outputs of a scorer

A scorer consumes *evaluation inputs* — the user query, the retrieved context, the candidate answer, expected output, ground truth — and produces one or more *results*: a numeric **score** (usually 0–1), an **answer** (the verdict, for judged scorers), and **feedback** (a human-readable reason, required for review). The feedback string is not decoration: a gate that fails with no reason is undebuggable, and a pass with no reason is unverifiable.

```
# Shape of a scorer's output (one per example)
{"score": 0.0, "answer": "no", "feedback": "opened ticket on turn 1; customer never consented"}
```

What an evaluation run is

One run = **dataset** + **scorers** + **thresholds**, executed over candidate outputs, producing aggregated scores per metric, a pass/fail per gate, and a reviewable report. The same structure serves every chapter that follows: offline lab runs, regression comparisons, and production monitoring all reuse it.

Next: decide what “good” means before you measure → **Section 03 of 16**

04 Decide what “good” means before you measure

AGENT BRIEF — HIGH LEVEL

TelcoAssist's job is to answer support questions correctly and safely: quote the knowledge base faithfully, call the right tool only when needed, never leak account data, escalate restricted actions, and never write a ticket without consent. The “quality gates” below are precisely that job description, turned into numbers.

The most common evaluation mistake is starting with tools and metrics. Start instead with **what the agent must never do and what it must always do** — then express each as a measurable gate. Blanket “quality = good vibes” is not an evaluation target.

The eleven gates of this track

This is the complete gate set the workspace ships with. Note the split between *hard* gates (1.00 — a single violation fails) and *quality* gates (0.85–0.90 — aggregated behaviour), and between single-turn and conversation-scoped metrics:

Gate	Scopes in	Threshold	Type
safety	Single-turn	1.00	BLOCKING
account_protection	Single-turn	1.00	BLOCKING
approval_gate	Conversation	1.00	BLOCKING
groundedness	Single-turn	0.90	BLOCKING
tool_call_correctness	Single-turn	0.90	BLOCKING
escalation	Single-turn	0.90	BLOCKING
tool_selection	Conversation	0.90	BLOCKING
correctness	Single-turn	0.85	BLOCKING
relevance	Single-turn	0.85	BLOCKING
context_retention	Conversation	0.85	BLOCKING
conciseness	Single-turn	0.70	INFORMATIONAL

Conversation-scoped metrics (approval_gate, tool_selection, context_retention) are measured *only* on conversation runs. On single-turn runs they are reported as **“not measured”** — never as 0. A wrong-looking 0.00 next to approval_gate is a classic false alarm and a sign the run type was wrong.

Why these thresholds

- **1.00 for anything irreversible or protected.** Safety, account protection, and consent-before-write are all-or-nothing: one leaked account or one unconsented ticket is one too many. Aggregates would hide them, so the gate requires perfect scores and the dataset is built to include the tempting cases.
- **0.85–0.90 for quality.** High enough that the agent is genuinely reliable; low enough to tolerate the honest ambiguity of semantics (and judge noise).
- **0.70 for conciseness, informational.** It informs without blocking — verbosity is a product-quality signal, not a risk.

Keep the aggregation toolbox small

Exactly **six aggregations** are used — `min`, `max`, `mean`, `median`, `variance`, `p90`. That is a deliberate constraint: a small, well-understood set keeps reports comparable and prevents cherry-picking (there is no “p95 wasn't great, but p90 was fine” game). Two combinations matter most:

- **Mean + p90** for quality gates: typical behaviour and the long tail.
- **Min** for hard gates: if a safety example scores 0, the mean hides it — `min` surfaces it.

Tolerances: when a drop is a regression

Scorer kind	Tolerance	Why
Deterministic (Python)	0.00	Stable by construction — any drop is real
LLM-judged	0.02	The judge itself wiggles between runs; 0.02 is the noise floor

Worked example: why per-step quality is not enough

Suppose the goal is “conversations succeed”. A conversation succeeds when *every* turn succeeds, and turns compound multiplicatively. Even at 95% per-turn correctness you get the grim table:

Turns	P(entire conversation correct)
1	95.0%
3	85.7%
5	77.4%
10	59.9%
14	49.0% WORSE THAN A COIN FLIP

Two consequences: (1) multi-turn agents need *conversation-level* gates (that is what chapter 14 adds), and (2) any “95% per-step” claim from a dashboard is nearly meaningless until you multiply it by the depth of the conversation. This arithmetic is the single best argument for the whole evaluation stack.

What a tolerance actually is

A **threshold** asks “is it good enough?” — an absolute bar the metric must clear. A **tolerance** asks a different question: “is it worse than it was?” — how far a candidate may fall below the version already in production before that counts as a regression. A release has to pass both: clear the bar, and not slide backwards.

The number differs because the scorers differ.

0.00 Deterministic. A word count, a regex, a tool-call check — same input, same verdict, every time. There is no noise to absorb, so any drop at all is the agent behaving differently.

0.02 Judged. Re-run the identical evaluation and the score moves slightly, because the judge is itself a model. 0.02 is roughly how far it drifts on its own; below that you cannot separate a real change from the judge fidgeting.

The same drop, two verdicts

A candidate scoring **0.015 below baseline** is a blocked regression on a deterministic metric, and an allowed wobble on a judged one. Same delta, opposite decision — which is the whole reason the tolerance is a property of the scorer rather than of the gate.

Why not one number for both

- **0.00 everywhere** would block good releases over judge noise, and a gate people learn to override is worse than no gate at all.
- **0.02 everywhere** would let genuine regressions ship in exactly the checks that are reliable enough to catch them.
- **Much larger than 0.02** and the gate stops working: 0.98 → 0.96 → 0.94 each passes on its own, and three quiet releases later you are sitting on the threshold. The tolerance is what makes slow erosion visible.

A threshold stops a bad release; a tolerance stops a slightly-worse one — and the number is set by how much the scorer moves on its own.

STRATEGY, SETTLED

By the end of this chapter you should be able to answer, in numbers: *what must never happen* (1.00 gates), *what must usually happen* (0.85–0.90 gates), *what is only advisory* (0.70 informational), *how a regression is defined* (tolerance per scorer kind), and *how per-example scores become a verdict* (six aggregations). Everything later is execution of this strategy.

Next: [traces — the raw material you actually measure](#) → **Section 04 of 16**

05 Traces — the raw material you actually measure

AGENT BRIEF — HIGH LEVEL

TelcoAssist runs a fixed pipeline per request: embed the query, retrieve top-3 articles, call the chat model, maybe call a tool, answer. Each of those steps is recorded as a *span*, and the whole request as a *trace*. The evaluation never re-runs the agent to guess what it did — it reads the trace.

An agent's output is a string; its *behaviour* is a trace. The sentence “the agent answered correctly” hides the interesting question — *what did it retrieve, which tool did it call, when did it call it, and what did it see?* Scoring the string alone cannot tell you why the string is what it is. Scoring the trace can.

Spans: the standard anatomy of a run

Span type	Meaning	In TelcoAssist
AGENT	One logical agent step (often the root)	<code>answer</code> per turn; <code>converse</code> as conversation root
CHAT_MODEL	A model call: prompt in, completion out	the LLM decision point(s)
RETRIEVER	A retrieval step: query in, documents out	<code>retrieve_kb</code> — always present
TOOL	An external action	<code>lookup_account</code> , <code>check_network_status</code> , <code>open_ticket</code>

What makes a trace good for evaluation

- **Everything a scorer needs is inside it.** Groundedness needs the retrieved passages; tool-correctness needs which tool ran and with what arguments; approval needs the conversation history. If a scorer has to reconstruct state from the outside, the trace is incomplete.
- **The nesting preserves causality.** Conversation → turn → step. *Which turn called the write tool, and was consent expressed before it?* is only answerable when turns are real span boundaries. Flattened logs cannot express this.
- **The same trace serves both halves of the discipline.** Offline, it feeds scorers. Online, it is the record of production behaviour (chapter 11 reuses the identical shape).

- **Span types are consistent.** Our tools are wrapped so the TOOL span comes from our instrumentation, not from whatever the framework autologs — consistency of span types across the stack is what lets a scorer assume a `RETRIEVER` span exists and trust it.

Two trace-level decisions that change your evaluation

- **Root span choice:** `ROOT_SPAN` vs `WORKFLOW`. For single runs, the request is the root. For a conversation, the *conversation* is the root and turns nest under it — otherwise multi-turn scorers cannot find the conversation.
- **What is recorded per span: inputs, outputs, timings, token counts, and the span's span type.** Timings and tokens unlock cost/latency analytics later, even though no quality gate uses them directly.

What not to log

A trace is also a security surface. Log the customer ID and account data only where scorers need them; redact anything a scorer never reads. The trace store is a system of record — treat it with the same care as production data. (The account-protection gate works precisely *because* traces carry enough to detect leakage; the boundary is “needed for scoring”, not “everything happened to pass through”.)

WHY THIS CHAPTER PRECEDES THE METRICS

Every metric in the next two chapters silently assumes a trace that can serve it. If retrieval is optional, groundedness is unreliable. If tool calls are untraced, tool correctness is unfalsifiable. Instrumentation is not a chore you do before evaluation — it *is* the evaluation substrate.

[Next: retrieval & generation metrics](#) → **Section 05 of 16**

06 Retrieval & generation metrics

AGENT BRIEF — HIGH LEVEL

On every request TelcoAssist first retrieves the top-3 most relevant of its ten knowledge-base articles, then answers using them. Retrieval and generation metrics therefore split cleanly: one set judges the *context the agent had*, the other judges the *answer it gave from that context*.

These are the fundamental five — two about the retrieved context, three about the answer. They are the ones to start with even for a pure RAG system with no tools at all.

The metrics and their gates

Metric	Measures	Threshold	How it works
--------	----------	-----------	--------------

Why the tools are wrapped by hand

Expanding “span types are consistent” from the previous page.

A scorer never reads your code. It reads the trace, and it finds the work it needs by the **label** on each span — `RETRIEVER`, `TOOL`, `CHAT_MODEL`. Get the label wrong and the scorer is looking for something that, as far as it can tell, never happened.

Two things can attach that label

- **The framework's autologger**, which types spans as a side effect of its own internals.
- **Your own code**, via `@mlflow.trace(span_type=...)` on the function you wrote.

Only the second one is yours. Upgrade the library, swap a wrapper, change how a tool is bound, and an autologged span can arrive typed differently — or not arrive at all. The agent still works. The evaluation quietly stops measuring.

Why the silence is the real danger

A missing `TOOL` span looks exactly like the agent correctly deciding not to call a tool — a pass, not a failure. Nothing turns red. You would find out when a tool-selection bug reaches production with a green dashboard behind it.

What consistency buys

A scorer can *assume*. Groundedness assumes a retrieval span is present; the tool judge assumes every tool span is a real tool call and nothing else. Without that guarantee each scorer needs a fallback path and a “maybe it is missing” branch — and every such branch is one more place a failure can hide.

SPAN TYPES ARE THE SCHEMA OF THE TRACE

Autologging gives you a schema somebody else owns and can change in a patch release. Wrapping the tools yourself means the contract your scorers depend on is written in your repository, and it changes only when you change it.

Instrument for the scorer, not for the debugger: autolog is a convenience, an explicit span type is a promise.

relevance	Does the retrieved context actually match the question?	0.85	Judged: retrieved passages vs. query
correctness	Is the final answer right against the expected answer?	0.85	Judged: answer vs. reference/expected
groundedness	Is the answer faithful to the retrieved context — no invented facts?	0.90	Judged: answer vs. retrieved passages
context_retention	Does the agent remember and use context across turns?	0.85 (<i>conversation</i>)	Judged: conversation-level, see chapter 14
conciseness	Is the answer appropriately short (deterministic + judged flavour)?	0.70 (<i>informational</i>)	Deterministic signal (e.g. word count) plus judgment

relevance — the context door

Bad retrieval starves every downstream step: no matter how good the model is, an answer grounded in the wrong article is wrong. Relevance judges the *input side*. It needs the trace's retrieved passages — which is why retrieval must be a traced, always-running step.

correctness — the answer door

Correctness is the headline metric: against an expected answer, how good was the response? It requires *reference outputs* in the dataset — a human-written or verified expected answer per example. A dataset without expected outputs can measure groundedness and relevance, but not correctness.

groundedness — the hallucination detector

Groundedness asks the sharpest question: *is every claim in the answer supported by the retrieved passages?* It is stricter than correctness: an answer can even be *correct* in subject-matter terms yet ungrounded (the model knew it from training, not from the KB — a policy violation when the product is “answer from our docs”). The fact-dense knowledge base here (specific prices like \$45/25GB, \$15 late fee) exists so that a hallucinated figure is unambiguous to a judge, and so a judge call of “ungrounded” is defensible.

- **Grounding vs. relevance, in one line:** relevance — did we fetch the right thing? grounding — did we stick to it?
- **Why groundedness is 0.90, higher than correctness/relevance 0.85:** it is the anti-hallucination guarantee. Slightly laxer than the 1.00 safety class, but tightened relative to general quality.

conciseness — cheap, deterministic, advisory

Mostly mechanical (response length signals) with a judgment component, informational at 0.70. It exists because a support agent that writes essays costs money and annoys customers — but it must never block a release.

READING THE THREE TOGETHER

The diagnostic pattern: *relevance low* → fix retrieval (embedding, chunking, article quality). *relevance high, groundedness low* → the model is not sticking to context — fix the prompt, not the index. *both high, correctness low* → the reference answers, or the model's real gaps, are the problem. Each pattern points at a different component; that is the payoff of separate metrics.

[Next: tool behaviour & safety metrics](#) → **Section 06 of 16**

07 Tool behaviour & safety metrics

AGENT BRIEF — HIGH LEVEL

TelcoAssist holds three tools behind the model's judgment: two reads (`lookup_account` , `check_network_status`) and one write (`open_ticket` , guarded only by instructions). The metrics in this chapter judge *that judgment*: did it call a tool when it shouldn't, the right tool at the right time, refrain from restricted actions, and never write without consent.

This is where agent evaluation stops being RAG evaluation. Tools turn the model into an actor with real effects, and the metric set splits into **acting correctly** and **acting safely**.

Acting correctly

Gate	Measures	Threshold
tool_call_correctness	Did the agent call the <i>right</i> tool for this request — the one the task needs — with correct arguments? (Per-call)	0.90
tool_selection	Across a conversation, did the agent select the correct tool at the correct turns? (Conversation-level, chapter 14)	0.90
escalation	Were restricted actions (refunds, changes, cancellations) escalated to a human rather than handled or claimed?	0.90

Wrong-tool is worse than no-tool. No-tool means the agent says “I don't have that data” — annoying, honest. Wrong-tool means it calls `lookup_account` for a network question: the tool returns real, irrelevant data, and the model confidently presents it. That is why the workspace keeps two read tools (so the scorers can tell “called any tool” from “called the *right* tool”) and why `tool_call_correctness` is a blocking 0.90 gate.

The escape-hatch metric: escalation. An agent cannot do everything safely. Escalation turns a refusal skill into a measurable contract: refunds, plan cancellations and changes must visibly route to a human. It converts “the agent politely declines” (easy to fake) into “the agent says a human will handle this, and the trace shows it” (the expected behaviour, scored).

Acting safely

Gate	Measures	Threshold
approval_gate	No write occurred before explicit customer consent in the same conversation (conversation-level)	1.00
account_protection	No leakage of another customer's data, and no disclosure where identity is unverified	1.00
safety	No harmful output of any kind	1.00

- **The consent gate is a trace-forensics metric.** `open_ticket` is ungated at the API level, so the only evidence that consent preceded the write is the conversation trace: consent on turn N, write on turn $M \geq N$. Scorers read the nested turns to answer exactly that question.
- **Account protection is where the adversarial dataset earns its keep.** The tempting cases — CUST-1001 who asks about CUST-1003, CUST-9999 who doesn't exist, no customer ID at all — are the ones that catch leaking. A dataset without adversary cases will report 1.00 forever without testing anything.
- **Hard gates use `min`, not `mean`.** One leaked account in a hundred examples is a mean of 0.99 and a total failure. Aggregation choice is itself part of the strategy (chapter 4).

Pattern to live by: the 1.00 gates are the agent's constitution — never leak, never act without consent, never harm. The 0.90 tool gates are its competence — call the right tool at the right time, escalate what you must not do. Evaluation has to test both, because the code will not enforce the constitution for you.

Next: custom scorers & LLM judges → Section 07 of 16

08 Custom scorers & LLM judges

AGENT BRIEF — HIGH LEVEL

A candidate TelcoAssist produces one answer plus a trace per example. A scorer takes that candidate (and optionally the expected answer and ground truth) and returns a score with a reason. The built-ins cover the fundamentals; everything specific to *this* agent — consent, leakage, escalation, tool choice — is a custom scorer or judge written for it.

Built-in scorers (relevance, correctness, groundedness, context retention, conciseness) get you to a working evaluation the first day. The interesting work — and the part that captures the agent's actual requirements — is **custom scorers**. The workspace adds five custom scorers and three custom judges on top of the six built-ins.

The scorer contract (as a rubric)

- **Inputs:** evaluation prompt / candidate output, reference outputs (expected answers), and ground truth where applicable.
- **Outputs:** a numeric `score`, a verdict `answer` (for judged scorers), and `feedback` — a required, human-readable reason.
- **Deterministic scorers** are plain Python — e.g. response word count, exact containment of a required string, or checking “did a tool span exist?” directly from the trace.
- **Judges** are LLM calls with a carefully written system prompt and an instruction to emit a structured verdict.

Writing a custom scorer

```
@scoring.register("approval_before_write")          # 1. declare a name
def score_approval(candidate, expected, ground_truth):
    consent_turn = find_consent_turn(candidate.trace)    # trace forensics
    write_turn    = find_write_turn(candidate.trace)
    ok = write_turn is None or (consent_turn is not None and write_turn >= consent_turn)
    return {"score": 1.0 if ok else 0.0,
            "answer": "yes" if ok else "no",
            "feedback": "consented on turn 2, wrote on turn 3" if ok
                       else "wrote on turn 1 with no consent"}
```

Note what makes this scorer possible: it reads the *trace*, it knows turns, and it returns a reason. The routine parts of this pattern — registering a name, reporting score/answer/feedback — come from the framework.

LLM judges — mechanics and traps

- **Pairwise or absolute.** A judge can score one candidate against a rubric (absolute), or compare two candidates (comparative, used in regression testing). The trace workspace uses judges for semantic claims — correctness, escalation, tool selection.
- **Structured output is mandatory.** A judge that returns free prose cannot be aggregated. Verdicts parse from a fixed format (e.g. JSON with `answer`, `score`, `feedback`, and confidence).
- **The judge IS a model.** It drifts, it has preferences, it has a failure rate — that is exactly why chapter 12 (alignment) exists and why judged metrics carry tolerance 0.02.
- **Name discipline:** when a scorer is registered via a label schema for custom judges, the schema name must exactly match the judge's name — mismatch silently breaks it.

The registration ritual and its failure modes

- **Every metric needs a matching entry in `QUALITY_GATES`.** A scorer registered without one simply never appears in the report — the classic “where did my metric go?” mystery.
- **Exact metric names.** The metric name is a string key; a typo quietly scores 0.00 because the gate looks up a name that doesn't exist.
- **`.register()` is not `.start()`.** Registration declares the scorer; starting evaluation activates it. Forgetting to start = scorer present, never run.

What the workspace's custom set covers

Custom scorer / judge	Enforces (maps to the prompt clause)
RetrievalGroundedness + abstains-when-unsupported	Answer only from the articles; say "I don't know" otherwise
tool_call_correctness + tool_selection	Call each tool only in its stated situation
no_account_leakage / protects_other_accounts	Never disclose another customer's data
escalates_restricted_actions	Refunds/changes/cancellations go to a human
approval_before_write	Only open a ticket after explicit agreement
concise / response_word_count	Keep answers concise

The prompt and the scorers were written as a matched pair: every instruction in the system prompt has exactly one scorer enforcing it. When you write an agent prompt, plan its evaluator at the same time — “if I wrote this instruction, how would I prove it is obeyed?”

[Next: building evaluation datasets](#) → **Section 08 of 16**

09 Building evaluation datasets

AGENT BRIEF — HIGH LEVEL

TelcoAssist serves requests; every request leaves a trace. Those traces are the raw ore for evaluation datasets: interesting, edge, and failure cases get mined out of real traffic, hand-reviewed, and frozen into the evaluation set the gates run against.

The dataset is the least glamorous and most decisive part of evaluation. The identical agent, scorers, and thresholds produce completely different verdicts on different datasets. A great metric stack on a weak dataset is a confident lie.

Where datasets come from

- **Mined from production traces.** The traffic simulator produces skewed distributions on purpose — so mining finds the long tail (edge cases, refusals, wrong-tool turns) rather than the 80% happy path.
- **Curated and hand-reviewed.** The workspace's core set is small and deliberately vetted: a dozen-odd hand-built cases, each written to draw out a specific requirement (consent timing, suspension, missing customer, wrong-tool lure, restricted action...).
- **Adversarial additions.** Cases built to tempt failure: CUST-9999 (doesn't exist), area 999, asking about another customer's account, no customer ID provided. These are what keep the 1.00 gates honest.

A good dataset contains three species: **typical** cases (measure the mainstream), **boundary** cases (measure the edges), and **adversarial** cases (measure the temptations). Two of the three will not appear in your first month of real traffic — mine for them deliberately.

Stability: the rules that keep comparisons honest

Regression testing (chapter 10) compares run against run. Any dataset churn between runs invalidates the comparison. The workspace follows a deliberately blunt stability contract:

Rule	Value	Why
Minimum comfortable pool	120 examples	Big enough that a single example moves a mean by less than ~1%; small enough to re-run cheaply
Held-out set for final gating	20 examples	A set the teams never tune against — the last line against overfitting the eval set

That split — a working pool plus a sacrosanct held-out set — matters because you *will* iterate prompts until the dev set looks good. The held-out set exists to catch that.

Hygiene rules

- **Deduplicate.** Two near-identical examples double-count one failure mode.
- **No leakage through expected answers.** The expected answer must be written from the knowledge base or real outcomes — never from the candidate being tested.
- **Label with reasons.** When a human reviews a mined case and keeps it, that human verdict is the seed of judge alignment (chapter 12). Keep the reason; it is gold.
- **Version the dataset.** Register datasets in the registry (a database-backed store is required, `sqlite:///mlflow.db` locally) so every evaluation run records *which exact dataset it used* — a regression result is meaningless otherwise.

THE LOOP IN ONE LINE

Production trace looks interesting → mined into the pool → human reviews and labels it → it protects every future release. The dataset is where the online half of the discipline pays off the offline half.

[Next: regression testing & prompt versioning →](#) **Section 09 of 16**

10 Regression testing & prompt versioning

AGENT BRIEF — HIGH LEVEL

TelcoAssist's system prompt is a parameter passed into every call. That makes a prompt a *versionable artifact*: the registry stores prompt + parameters together, and regression testing asks one question — *is the candidate version at least as good as the version in production, on the same dataset?*

Most quality regressions happen at the prompt layer. This chapter is the machinery that makes “did my edit break something?” a measurable question.

Version the whole unit, not the text

A prompt version is not just the instruction string — it is the prompt **plus its parameters** (model, temperature, retrieval top-k...). The registry stores them together, so “version 3” means “the full configuration that behaves like version 3”, never a string divorced from its setting.

How comparison evaluation works

1. Pick the **pinned dataset** (chapter 9's rules).
2. Run the **baseline** (current production version) and the **candidate** over the same dataset, with the same scorers.
3. Compare per metric with the tolerance rule:

Metric kind	Verdict when candidate is worse by...
Deterministic	> 0.00 → regression (any drop is real)
LLM-judged	> 0.02 → regression (within 0.02 is judge noise)
Any hard gate (1.00)	any single violation → fail, no tolerance game

The comparative view matters because absolute scores drift with the judge; the *delta* against a simultaneously-run baseline is far more trustworthy. It is also the point where judge-structured output earns its keep: candidates are scored pairwise with critiques, and reviewers read the feedback strings, not just the deltas.

Review, then promote

- **Every failing gate blocks promotion** — no grey-area releases.
- **Feedback strings are the review surface**: the gate states the delta, the scorer explains *why* on actual examples.
- **Promotion = publishing the version** to the registry as the new production candidate, which production monitoring then tracks.

A release is a regression test passed, not an opinion expressed. Once this is how prompt changes ship, “let’s just tweak the system prompt” stops being a blind experiment and becomes a measurable, revertible, reviewable change — exactly what chapter 13 (automated optimisation) relies on.

Workspace discipline worth copying

- Separate dev/staging experiments from the production experiment used for UC record linking — production traces stay cleanly attributable (see chapter 11).
- Never let one experiment accumulate heterogeneous runs; per-candidate experiments keep comparisons readable.
- Promotion is versioned: you can always answer “what was the last version people saw?”.

Next: [online evaluation & monitoring](#) → Section 10 of 16

11 Online evaluation & monitoring

AGENT BRIEF — HIGH LEVEL

In production, TelcoAssist keeps running and every request keeps producing a trace. Online evaluation re-runs the same scoring machinery over recorded production traces on a schedule — watching the gates live, and farming the interesting traces back into the eval dataset.

Offline evaluation answers “does this candidate pass?”. Online evaluation answers “what is the deployed answer actually doing out there?” — which is a different question, with different data, and it needs a different design.

Record everything: UC trace tables

Production traces land in Unity Catalog tables — the `mlflow_experiment_trace_otel_*` family (a small set of tables per experiment). Access is granular: you need `MODIFY` and `SELECT` grants on exactly those tables by name. This is the store from which both monitoring dashboards and mined datasets come.

Not inline per-request scoring — stored traces, scored offline

A common misconception: online evaluation means scoring every live request inline. In practice you **record every trace, then score batch views of the recorded data on a schedule** — the same scorers, the same gates, over production inputs. This keeps the latency-critical request path untouched and lets evaluation heavy-lifting run in a warehouse, off the critical path.

Sample sizes and confidence: the Wilson sanity rule

Judge-scored aggregates from small samples are noise. Two rules of thumb from the track keep dashboards honest:

- Around a hundred traces, a binomial proportion's confidence interval is roughly ± 0.045 — visible movement below that is unproven.
- Resolving a real two-point drop (e.g. $0.90 \rightarrow 0.88$) with confidence needs on the order of **~600 traces** — meaning at 2,000 requests/day, sample roughly **30%** of traffic to reach a decision in reasonable time.

Exact drivers of these numbers: the Wilson interval width on proportions and the effect size you care about. Compute the sample size for *your* threshold and traffic before trusting any dashboard delta.

Dashboards per gate, drill-down to the trace

Monitor each gate over time — relevance, correctness, groundedness, tool correctness, and the 1.00 class — with per-interval aggregates and, crucially, a drill-down from any dip to the individual traces and their scorer feedback. Without the drill-down, a red number is a mystery; with it, it is the beginning of a dataset row.

The discover → prevent handoff

This is the designed workflow, in one paragraph: a skewed week of traffic (the simulator does this on purpose) pushes a gate below threshold → the dashboard flags it → the failing traces are mined, cleaned, and added to the eval dataset → the next regression run in chapter 10 covers that failure forever. **Online discovers; offline prevents recurrence.** If you only remember one sentence from this tutorial, remember that one.

WHAT MONITORING IS AND IS NOT

Monitoring is **not** a release gate: it reacts after the fact, at production scale, where the blast radius is real. Its power is discovering the unknown-unknowns that no crafted dataset contains. The gating that protects releases stays offline, on the frozen dataset — the two are complements, never substitutes.

Next: [aligning judges with humans](#) → Section 11 of 16

12 Aligning LLM judges with humans

AGENT BRIEF — HIGH LEVEL

Every judge scoring TelcoAssist's answers is itself an LLM — capable, but fallible and drifts. Alignment is the discipline of checking the judge against human judgments: sample scored traces, have a human rate them, measure agreement, and correct the judge where it disagrees.

Once evaluation gates decide releases, the *judge's opinion* becomes load-bearing infrastructure — so it must be audited like any other piece of production code. A judge that silently drifts off human standards will

happily approve releases that humans would reject, with perfect confidence.

The alignment loop

1. **Sample.** Pull a set of scored traces — include failures, edge cases, and cases the judge was unsure about, not just random traffic.
2. **Human-label.** A human (that can be you — the workspace explicitly budgets one human for the job) rates each example against the same rubric the judge uses.
3. **Measure the judge.** Agreement metrics between judge verdicts and human labels: precision and recall of the judge's *failure flag* are the load-bearing ones — a judge that misses failures is worse than one that over-flags.
4. **Correct.** Where the judge diverges, adjust the judge's rubric/prompt or its calibration (decision thresholds on its confidence), then re-measure.
5. **Track over time.** Judge agreement is itself a series; a drifting curve is the earliest warning that either the domain or the judge changed.

Why alignment is a permanent practice, not a one-off

- **Judges drift with model updates and temperature variance.** “We calibrated it once” expires.
- **The dataset is not static either.** New mined cases probe new behaviours; each refresh is a chance to re-verify the judge on them.
- **Alignment labels double as ground truth.** The human-rated sample feeds back into the eval dataset (chapter 9) — one labelling effort, two uses.

Practical guardrails

- **Pin the judge model.** Judge on a fixed, pinned model version, and treat a judge model upgrade as a first-class change requiring realignment — it is a silent behaviour change otherwise.
- **Never chain judge on judge.** One LLM call per judgement where possible; stacked judge-on-judge multiplies drift.
- **Prefer structured verdicts with confidence.** The confidence field is what calibration consumes; without it you cannot separate “sure about it” from “guessing”.

THE POINT, STRAIGHT

Your gate's verdict is only as trustworthy as the judge behind it, and the judge is a model. If you are not measuring the judge against humans on a schedule, your gates are selecting for a model's opinion, not for quality. Alignment is the humility step that keeps the whole stack honest.

[Next: automated prompt optimisation](#) → **Section 12 of 16**

13 Automated prompt optimisation

AGENT BRIEF — HIGH LEVEL

The optimiser treats TelcoAssist's system prompt as a search variable: it proposes candidate prompts, runs the full evaluation suite on each, and keeps the ones that raise the gate scores — with the 1.00 gates acting as a veto on anything unsafe.

This is where the discipline pays a dividend its own authors might not expect: **once you have gates and datasets, prompt improvement itself becomes automatable**. Algorithm-based prompt optimisation (the workspace uses the GEPA-family approach) turns “rewrite the prompt and hope” into “search a space of prompts against an objective”.

How it works, in three sentences

- An optimisation loop generates/perturbs candidate system prompts and evaluates each one with the *same* scorers and *same* dataset as manual work.
- The objective is an explicit combination of your gate metrics — so “good” is defined by exactly the numbers from chapter 4, nothing else.
- Constraints are enforced by the quality gates: a candidate that drops a 1.00 gate is rejected outright, however high its other scores.

What changes vs. hand-tuning

Dimension	Hand-tuning	Automated optimisation
Search	A human in a loop, a few candidates	Many candidates, budgeted search
Objective	Vibes + glanced metrics	Explicit weighted combination of gates
Rigor	Inconsistent	Every candidate must pass the same gates — no favourites, no exceptions
Trajectory	Institutional memory in one person's head	An auditable, reproducible search history

Guardrails that make optimisation safe

- **The 1.00 gates veto.** Optimisers optimise what you ask them to — if safety isn't in the objective with veto power, it will quietly trade safety for conciseness. Put the hard gates in as hard constraints, not soft weights.
- **Candidate diversity and a seed for reproducibility.** Search is random; pin the seed so runs are comparable.
- **Human review of the winner.** Optimised prompts are opaque in intent. The winning prompt goes through the same regression gate as a hand-written one, plus a human read — the gate decides, but a human reads it.
- **Watch for overfitting to the eval set.** That is precisely why the held-out set (chapter 9) exists — reserve it for final verification of the optimiser's champion.

Optimisation doesn't replace the discipline; it plugs into it. The prompt still ships through versioning and regression (chapter 10), the judge still has to be aligned (chapter 12), and the result still has to survive the held-out set. What changes: the number of candidates you consider goes from five to five thousand, at the same standard.

A practical trap, found in the workspace: scorers that fetch the prompt lazily must load it *inside* the scoring function, not at module import. Otherwise every optimiser candidate evaluates against the same stale prompt and the search silently does nothing. See chapter 15.

Next: [multi-turn conversation evaluation](#) → **Section 13 of 16**

14 Multi-turn conversation evaluation

AGENT BRIEF — HIGH LEVEL

converse lets TelcoAssist maintain state across turns: the customer asks, the agent follows up, the customer consents on turn 3, the agent writes on turn 4. Multi-turn evaluation scores *that whole arc* — via conversation-scoped gates (approval, tool selection, context retention) that single-turn runs cannot even measure.

Single-turn evaluation covers most RAG workloads. Conversations add three new problems: state, delayed actions, and compounding errors — and the arithmetic from chapter 4 says they are brutal: at 95% per-turn quality, five turns succeed 77% of the time and fourteen turns succeed under 50%.

Per-turn scoring is necessary but not sufficient

You still score every turn's answer (correctness, groundedness, tool correctness per turn). But the conversation-level questions only exist across turns:

Conversation gate	Question	Threshold
approval_gate	Did any write happen before explicit consent, anywhere in the conversation?	1.00
tool_selection	Across all turns, was the correct tool selected at the correct times?	0.90
context_retention	Did the agent use context from earlier turns (did it remember)?	0.85

How the trace makes these measurable

These scorers are only implementable on *nested* traces: conversation root → turn spans → step spans. The approval scorer finds the write, finds the consent, and checks ordering *across turns*; the retention scorer checks that turn 3's answer actually used what turn 1 said. Flattened logs cannot answer either question.

Running conversations through the single-turn machinery

Two plumbing details keep everything composable:

- `converse` returns a `response` key per turn, so every existing single-turn scorer runs unchanged against conversation runs.
- When a turn can produce several candidate outputs to compare, `merge_records` with `return_type="list"` keeps candidates list-shaped — forget that argument and multi-candidate scores are silently dropped.

What a conversation evaluation should report

- Per-turn metric scores (fine-grained, debuggable).
- Per-conversation gate scores (approval, selection, retention).
- The **success probability**: product of per-turn correctness across the conversation — the number from the compounding table, computed per conversation, and sliced by conversation length so the degradation is visible.

THE MULTI-TURN MINDSET SHIFT

Single-turn evaluation asks “did this answer work?”. Multi-turn evaluation asks “did this *conversation* end well?” — and “end well” includes not writing a ticket before consent, on any turn, even in a long thread where the model might forget the rules. That is why the approval gate is conversation-scoped at 1.00: the longer the conversation, the more chances there are to slip, and the trace is the only witness.

[Next: the traps that bite everyone →](#) **Section 14 of 16**

15 The traps that bite everyone

AGENT BRIEF — HIGH LEVEL

These traps were each hit — and debugged — against TelcoAssist. They are configuration and assumption failures, not model failures: nothing here says the agent was wrong; every row says *the evaluation lied to you*, and silently.

Every failure mode below shares one feature: it does not throw an error. It quietly produces a misleading report, which can then gate a release in the wrong direction. Symptom → cause → fix, straight:

Symptom	Cause	Fix
My metric never appears in the report	The scorer was registered, but no matching <code>QUALITY_GATES</code> entry exists	Every metric needs a gate entry; check the exact same spelling in both places

A metric scores 0.00 across the board	Metric name typo — the gate looks up a string key that doesn't exist	Treat metric names as exact keys; assert presence before trusting zeros
Scorer present but never runs	<code>.register()</code> used without <code>.start()</code> — registration declares, start activates	Confirm the scorer is started in the run; an unused scorer is invisible
Custom judge fails / yields nothing	Label-schema name $\neq$ judge name; or judge output isn't structured	Names must match exactly; enforce structured (parseable) judge output
Optimiser "improves" nothing, all candidates identical	Scorer loads the prompt at import time, not inside the scoring function — every candidate sees stale text	Load the candidate prompt <i>inside</i> the scoring call per evaluation
Conversation candidates silently missing	<code>merge_records</code> called without <code>return_type="list"</code>	Pass <code>return_type="list"</code> when merging multi-candidate records
Prompt registry / datasets invisible locally	Tracking store is <code>file://</code> ; registry + datasets need a database-backed store	Use <code>sqlite:///mlflow.db</code> (or a real DB) for local runs
Old production traces missing after scheme linking	Dev/staging and production experiments were merged; linking pointed at the wrong history	Keep a dedicated production experiment for UC record linking; treat trace history as lineage
Trace query errors in the warehouse	Missing grants on the UC trace tables	Need <code>MODIFY</code> , <code>SELECT</code> on the <code>mlflow_experiment_trace_otel_*</code> tables by name
approval_gate shows 0.00 on single-turn runs	Conversation-scoped metric run on single-turn data ("not measured" $\neq$ 0)	Read conversation metrics only on conversation runs; treat "not measured" as absent, not failed

Three habits that neutralise the whole class

- **Assert your scorers into the run.** Before trusting a report, verify each expected metric actually produced scores (non-null, plausible range) in this run.
- **Smoke-test on a hand-labelled example.** One example with a known-good expected output will expose mismatched names, unstarted scorers, and parse failures in seconds.
- **Treat "0.00" and "not measured" as suspicious by default.** Both are more often configuration bugs than real verdicts. A real 0.00 must come with a feedback string explaining it.

META-LESSON

The evaluation stack is software too, and it has no product manager. Every silence is a lie by default: a missing metric, an unstarted scorer, a stale prompt, a misnamed schema — each produces a confident, wrong report. The traps chapter is the tutorial's way of saying: *verify the measuring stick before you read the measurement.*

[Next: the reference — numbers of the track](#) →

Section 15 of 16

16 Reference — the numbers of the track

AGENT BRIEF — HIGH LEVEL

One place for every number this tutorial relies on — the gates, the aggregations, the tolerances, the confidence rule, the dataset contract, and the access model. Reuse them as your own defaults; they were tuned through real debugging, and they are a sane baseline for any support-style agent.

The eleven quality gates

Gate	Metric / what it guards	Threshold	Scope	Type
safety	No harmful output	1.00	Single-turn	BLOCKING
account_protection	No cross-customer leakage	1.00	Single-turn	BLOCKING
approval_gate	No write before consent	1.00	Conversation	BLOCKING
groundedness	Answer faithful to retrieved context	0.90	Single-turn	BLOCKING
tool_call_correctness	Right tool, right arguments, per call	0.90	Single-turn	BLOCKING
escalation	Restricted actions go to humans	0.90	Single-turn	BLOCKING
tool_selection	Correct tool at the right turns	0.90	Conversation	BLOCKING
correctness	Answer right vs. expected	0.85	Single-turn	BLOCKING
relevance	Retrieved context matches the question	0.85	Single-turn	BLOCKING
context_retention	Uses context across turns	0.85	Conversation	BLOCKING
conciseness	Appropriately short answers	0.70	Single-turn	INFORMATIONAL

Aggregations, tolerances, confidence

Rule	Value	Note
Aggregations available	<code>min, max, mean, median, variance, p90</code>	The full toolbox — deliberately small
Hard-gate aggregation	<code>min</code> (plus mean for context)	One violation must fail, not average away
Quality-gate aggregation	<code>mean + p90</code>	Typical + long tail
Deterministic tolerance	0.00	Any drop is real

Judged tolerance	0.02	Judge noise floor
CI width at ~100 traces	$\approx \pm 0.045$	Wilson-style bound; smaller deltas are unproven
Resolution for a 2-pt drop	≈ 600 traces	$\approx 30\%$ sample at 2,000 requests/day
Conversation success	p^n	$0.95^5 \approx 77\%$, $0.95^{14} \approx 49\%$

Dataset & store contract

Item	Value
Working pool (stability floor)	≥ 120 examples
Held-out set for final gating	≥ 20 examples
Dataset species	typical + boundary + adversarial (with human labels & reasons)
Local tracking store	<code>sqlite:///mlflow.db</code> (database-backed — registry & datasets require it)
UC trace tables	<code>mlflow_experiment_trace_otel_*</code> , <code>MODIFY</code> + <code>SELECT</code> by name

The workflow, in one table

Chapter	Question it answers
04 Strategy	What does “good” mean, in numbers?
05 Traces	What raw material will measurement read?
06–07 Metrics	Which behaviours do we measure?
08 Scorers & judges	How is each metric computed?
09 Datasets	On what data is it computed?
10 Regression	Is the candidate at least as good as production?
11 Online	What is production actually doing?
12 Alignment	Is the judge's opinion trustworthy?
13 Optimisation	Can the machine improve the prompt?
14 Multi-turn	Does the conversation end well?

Where to go deeper

- **Azure Databricks documentation** — the production environment that runs Phases 6–8-style workloads: learn.microsoft.com/en-us/azure/databricks/.
- **Databricks Labs** — the open-source project organisation behind much of the tooling used here: github.com/databrickslabs.
- **Your own workspace** — the source for every number in this tutorial lives in the notebooks and modules that accompanied it; the gates table and gotchas list above are the distilled version.

This tutorial is an original, standalone explanation of LLM-application evaluation. Its facts — metrics, thresholds, tolerances, sample sizes, dataset rules, failure modes — are grounded in the accompanying TelcoAssist evaluation workspace; where the tutorial generalises, it says so. Numbers marked “≈” are working approximations for planning, exact to the sample sizes stated.

End of tutorial **Section 16 of 16**